

---

# Complexity and Advanced Algorithms

## Module 2

### Parallel Computing

---

# Why Parallel Computing?

---

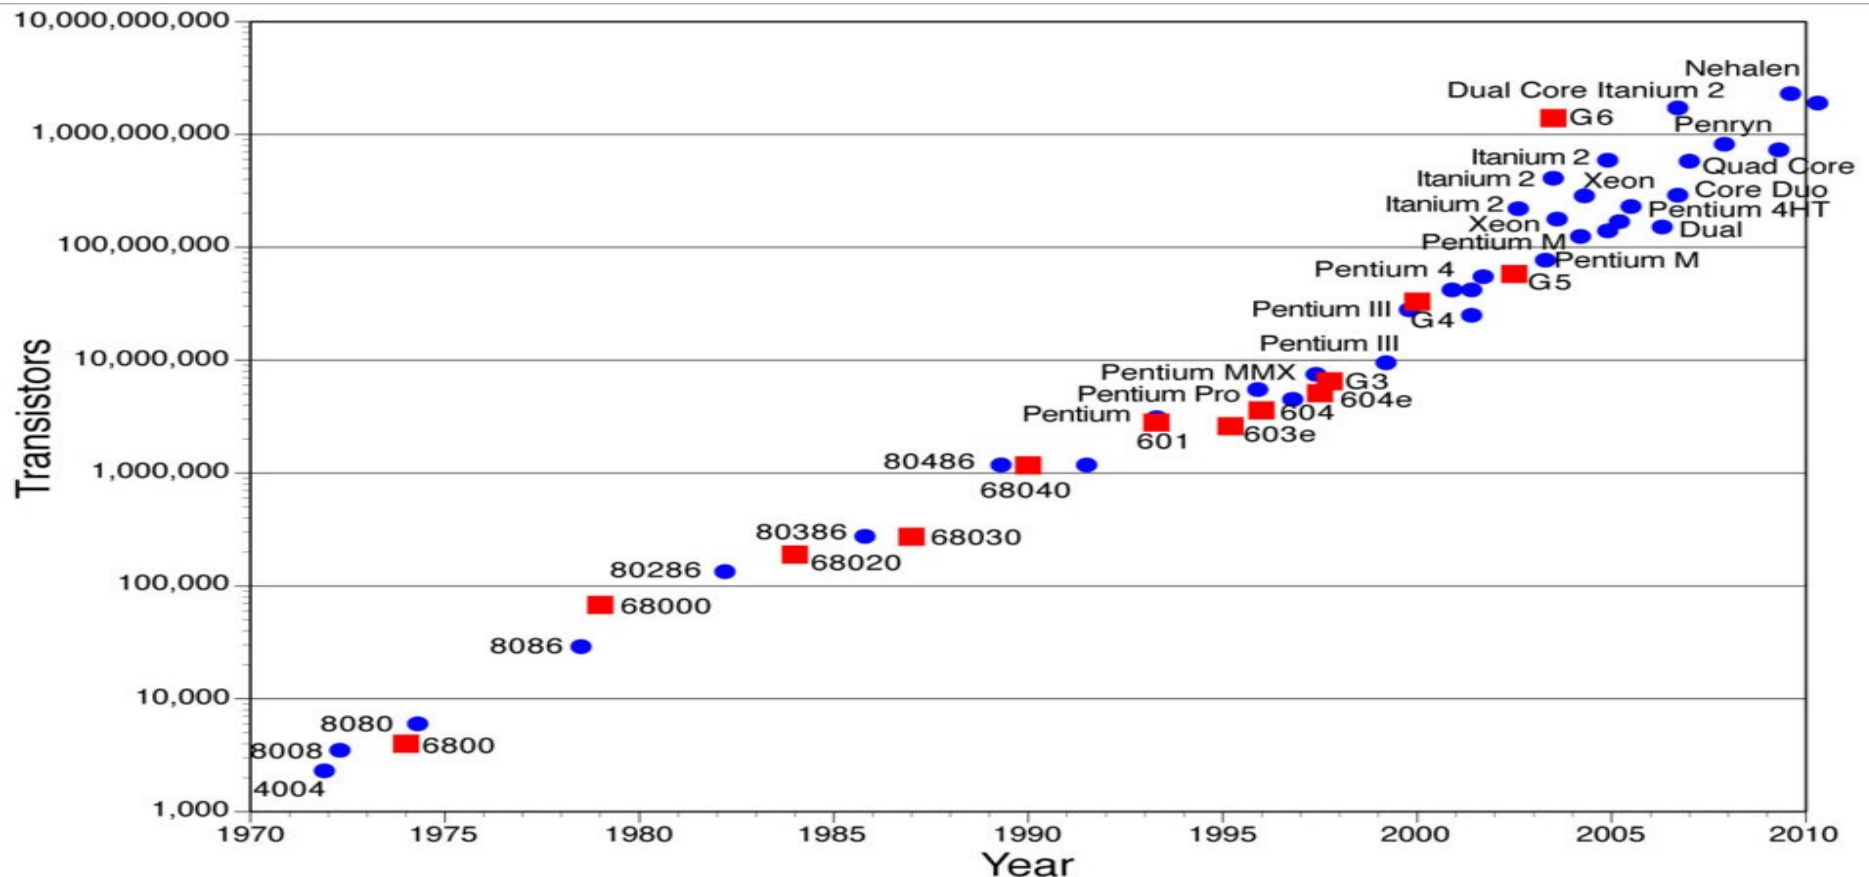
- Save time, resources, memory, ...
- Who is using it?
  - Academia
  - Industry
  - Government
  - Individuals?
- Two practical motivations:
  - Application requirements
  - Architectural concerns.
- Why now?
  - Most computers including laptops are multi-core!
  - Need to therefore study how to use parallel computers.

# 1. Application Requirements

---

- Several applications are pushing the limits with huge compute requirements:
  - Deep learning
  - Image/Video/Text search, retrieval, and indexing
  - Digital effects/computer graphics/animation
  - Materials/Life Sciences/Drug design/...
  - Social computing/web/...

## 2. Architectural Advances



- Moore's Law: The number of transistors that can be inexpensively placed on an integrated circuit is increasing exponentially, doubling approximately every two years.

# On the Other Hand...

---

- Present Difficulties
  - Memory Wall
  - Power Wall
  - ILP Wall

# The Brick Wall - 1

---

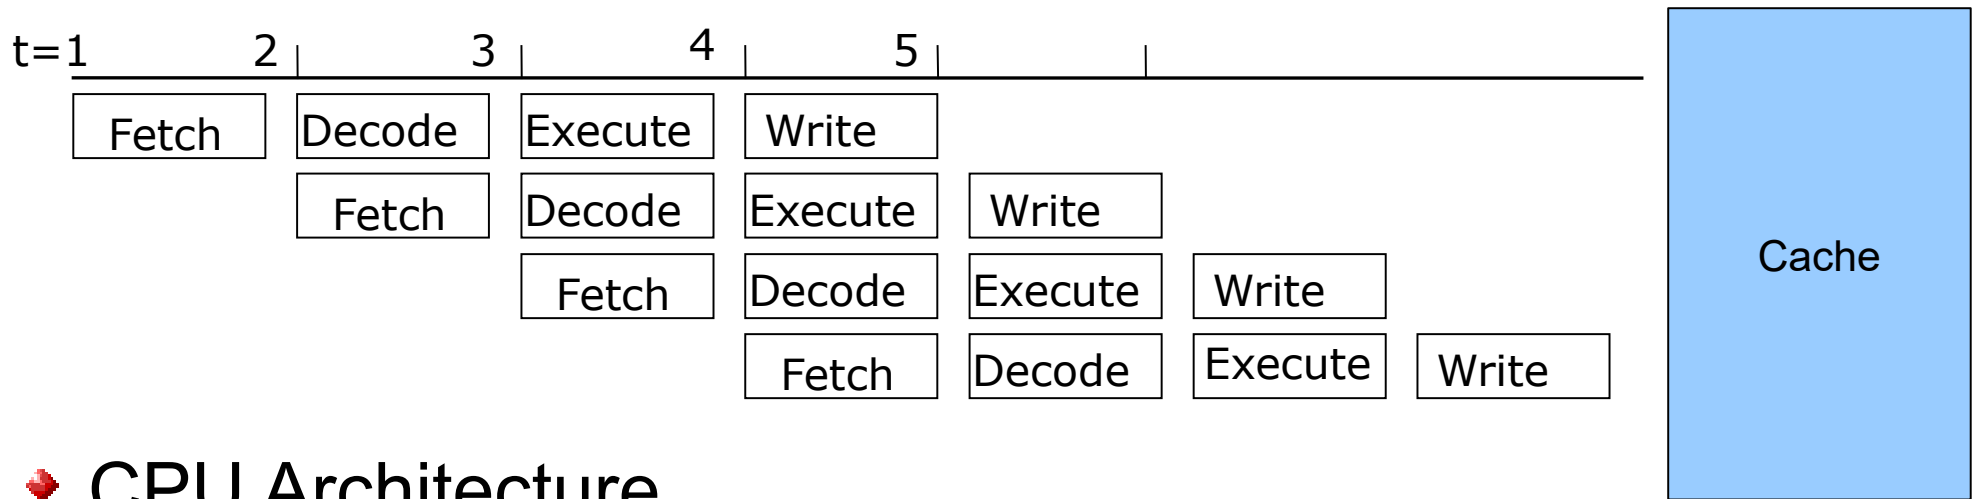
- Memory Wall
  - Memory latency up to 200 cycles per load/store.
  - Floating point operations take no more than 4 cycles.
  - Earlier, it was thought that “multiply is slow but load and store is fast”.

# The Brick Wall - 2

---

- Power Wall
  - Enormous increase in power consumption.
  - Power leakage.
  - However, presently “Power is expensive but transistors are free”.

# Basic Architecture Concepts

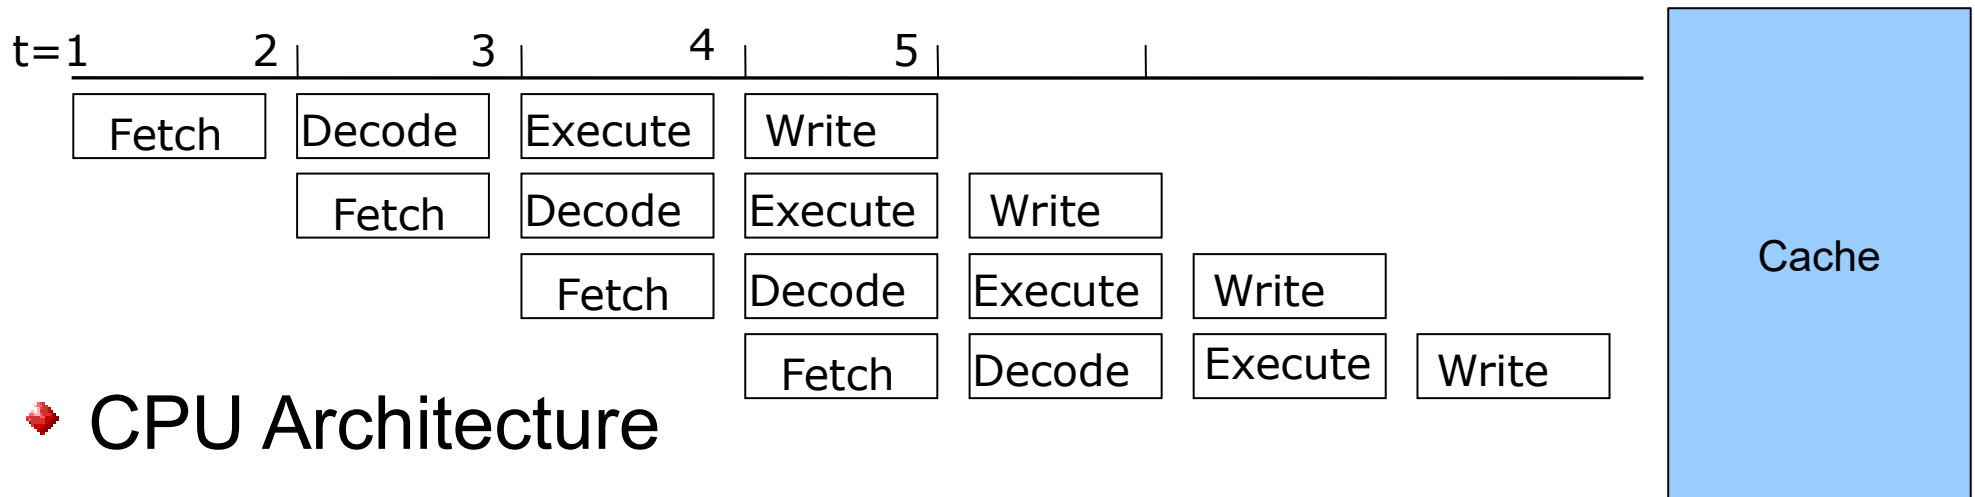


## ❖ CPU Architecture

- 4 stages of instruction execution
  - ▶ Too many cycles per instruction (CPI)
- To reduce the CPI, introduce pipelined execution
  - Needs buffers to store results across stages.
  - ▶ A cache to handle slow memory access times



# Basic Architecture Concepts



## ❖ CPU Architecture

- 4 stages of instruction execution
  - ▶ Too many cycles per instruction (CPI)
- To reduce the CPI, introduce pipelined execution
  - Needs buffers to store results across stages.
  - ▶ A cache to handle slow memory access times
  - Caches, out-of-order execution, branch prediction, ...

# The Brick Wall - 3

---

- ILP Wall
  - ILP via branch prediction, out-of-order and speculative execution
  - Diminishing returns from instruction level parallelism.

# Conventional Wisdom in Computer Architecture

---

- Power Wall + Memory Wall + ILP Wall = Brick Wall
- Old CW: Uniprocessor performance 2X / 1.5 yrs
- New CW: Uniprocessor performance only 2X / 5 yrs?

# Multicores to the Rescue

---

- Predicted that 100+ core computers would be a reality soon.
- Increased number of cores without significant improvement in clock rates.
  - Due to silicon technology improvements
- Big questions
  - How to exploit these cores in parallel?
  - What are the killer applications that can democratize these new models?
    - Search, web, ???

# The Academic Interest

---

- Algorithmics and complexity
  - How to design parallel algorithms?
  - What are good theoretical models for parallel computing?
  - How to analyze parallel algorithms?
  - Can every sequential algorithm be parallelized?
  - What are some complexity classes wrt parallel computing?

# The Academic Interest

---

- Systems and Programming
  - How to write parallel programs?
  - What are some tools and environments.
  - How to convert algorithms to efficient implementations.
  - What are the differences to sequential programming?
  - What are the performance measures?
  - Can sequential programs be automatically converted to parallel programs?

# The Academic Interest

---

- Architectures
  - What are standard architectural designs?
  - What new issues are raised due to multiple cores?
- Downstream concerns
  - Does a programmer have to worry about this?
  - How to support the systems software as architecture changes?

# The Course Coverage


---

- Focus on algorithms and complexity
- Models for parallel algorithms
- Algorithm design methodologies with application
  - Semi-numerical
  - Lists
  - Trees and graphs
- Some parallel programming practice
- Complexity, characterization, and connection to sequential complexity classes.



# The PRAM Model

---



- An extension of the von Neumann model.

# The PRAM Model

---

- A set of  $n$  identical processors
- A common access shared memory
- Synchronous time steps
- Access to the shared memory costs the same as a unit of computation.
- Different models to provide semantics for concurrent access to the shared memory
  - EREW, CREW, CRCW(Common, Arbitrary, Priority, ...)

# The Semantics

---

- In all cases, it is the programmer to ensure that his program meets the required semantics.
- EREW : Exclusive Read, Exclusive Write
  - No scope for memory contention.
  - Usually the weakest model, and hence algorithm design is tough.
- CREW : Concurrent Read, Exclusive Write
  - Allow processors to read simultaneously from the same memory location at the same instant.
  - Can be made practically feasible with additional hardware

# The Semantics

---

- **CRCW : Concurrent Read, Concurrent Write**

- Allow processors to read/write simultaneously from/to the same memory location at the same instant.
- Requires further specification of semantics for concurrent write. Popular variants include
  - **COMMON** : Concurrent write is allowed so long as the all the values being attempted are equal. Example: Consider finding the Boolean OR of n bits.
  - **ARBITRARY** : In case of a concurrent write, it is guaranteed that some processor succeeds and its write takes effect.
  - **PRIORITY** : Assumes that processors have numbers that can be used to decide which write succeeds.

# PRAM Model – Advantages and Drawbacks

---

## Advantages

- A simple model for algorithm design
- Hides architectural details for the designer.
- A good starting point

## Disadvantages

- Ignores architectural features such as:
  - memory bandwidth,
  - communication cost and latency,
  - scheduling, ...
- Hardware may be difficult to realize


# Example 1 – Matrix Multiplication

---

- One of the fundamental parallel processing tasks.
- Applications to several important problems in linear algebra, signal processing and optimization.
- Several techniques that work in parallel also.

# Example I – Matrix Multiplication

- Recall that in  $C = A \times B$ ,  $C[i,j] = \sum A[i,k].B[k,j]$ .
- Consider the following recursive approach:
  - Works well in practice.



$$C_{00} = A_{00} \cdot B_{00} + A_{01} \cdot B_{10}$$

$$C_{01} = A_{00} \cdot B_{01} + A_{01} \cdot B_{11}$$


$$C_{10} = A_{10} \cdot B_{00} + A_{11} \cdot B_{10}$$

$$C_{11} = A_{10} \cdot B_{01} + A_{11} \cdot B_{11}$$

# Example I – Matrix Multiplication

---

- Other approaches include Cannon's algorithm



- Can overlap computation with communication.
- Works well when the number of processors is more.



# Example 2 – New Parallel Algorithm

---

Listing 1:

$S(1) = A(1)$

for  $i = 2$  to  $n$  do

$S(i) = S(i-1) \circ A(i)$

- **Prefix Computations:** Given an array  $A$  of  $n$  elements and an associative operation  $\circ$ , compute  $A(1) \circ A(2) \circ \dots \circ A(i)$  for each  $i$ .
- A very simple sequential algorithm exists for this problem.

# Parallel Prefix Computation

---

- The sequential algorithm in Listing 1 is not efficient in parallel.
- Need **a new algorithm approach**.
  - **Balanced Binary Tree**

# Balanced Binary Tree

---

- An **algorithm design approach** for parallel algorithms
- Many problems can be solved with this design technique.
- Easily amenable to parallelization and analysis.


# Balanced Binary Tree

---

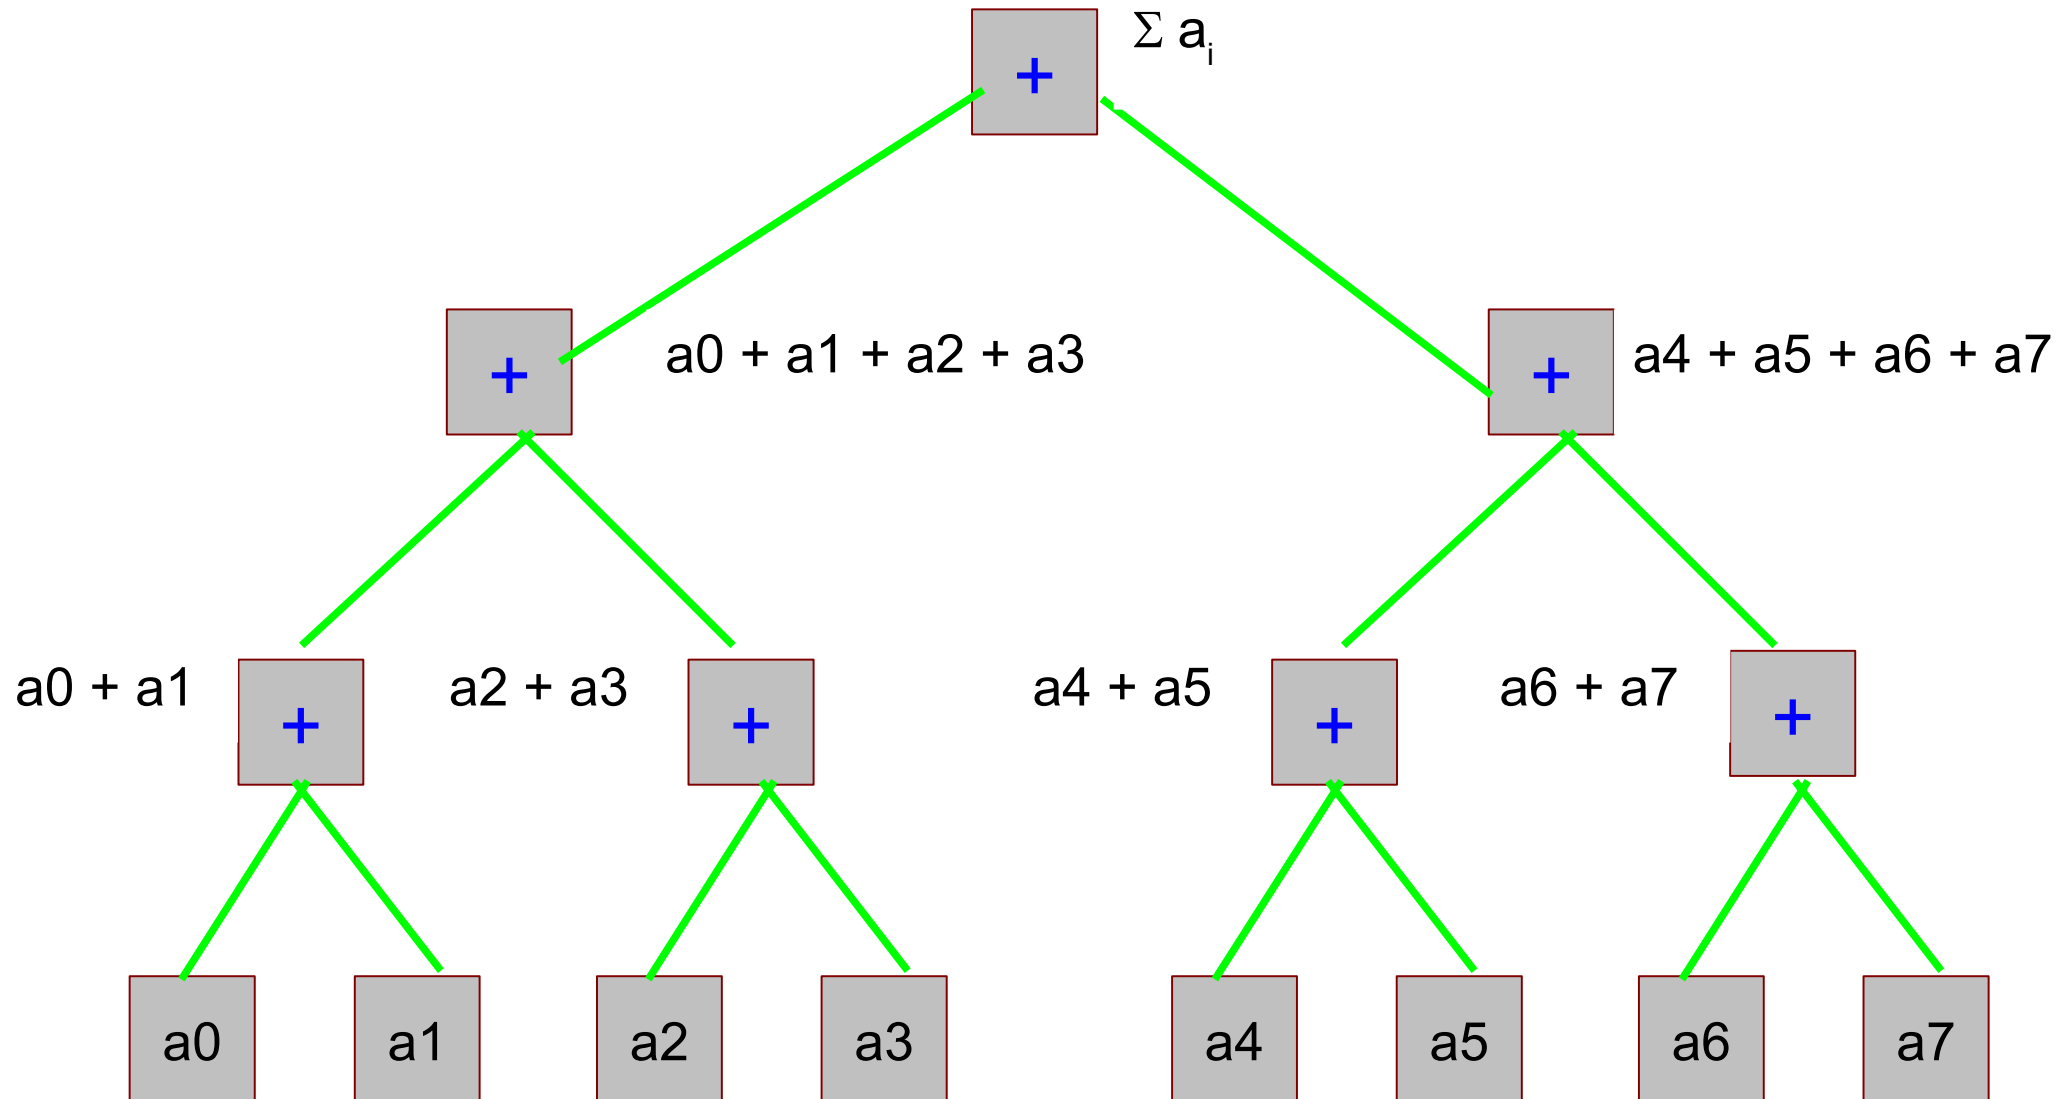
- A complete binary tree with processors at each internal node.
- Input is at the leaf nodes
- Define operations to be executed at the internal nodes.
  - Inputs for this operation at a node are the values at the children of this node.
- Computation as a tree traversal from leaf to root.

# Balanced Binary Tree – Prefix Sums

---



# Balanced Binary Tree – Sum



# Balanced Binary Tree – Sum

---

- The above approach called as an "upward traversal"
  - Data flow from the children to the root.
  - Helpful in other situations also such as computing the max, expression evaluation.
- Analogously, can define a downward traversal
  - Data flows from root to leaf
  - Helps in settings such as element broadcast

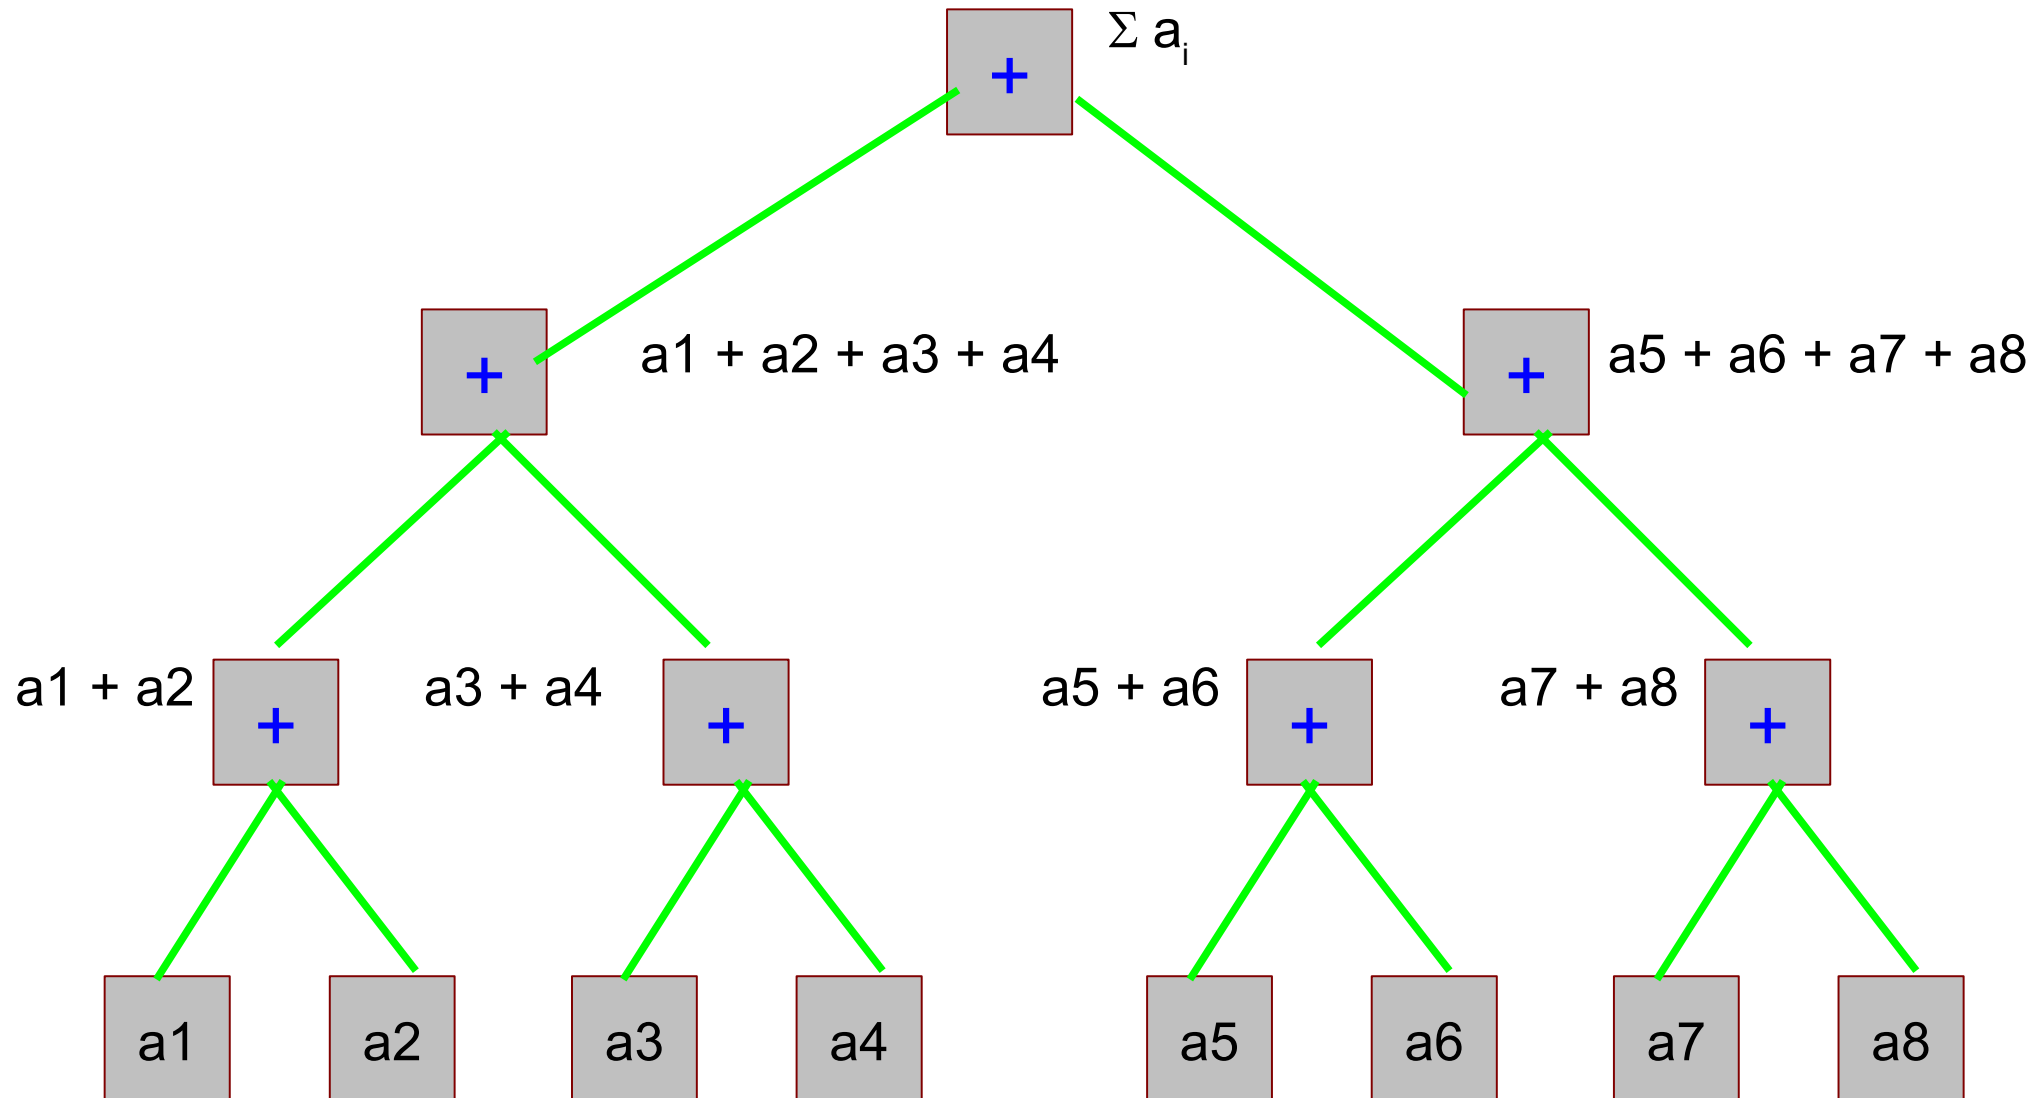
# Balanced Binary Tree

---

- Can use a combination of both upward and downward traversal.
- Prefix computation requires that.
- Illustration in the next slide.

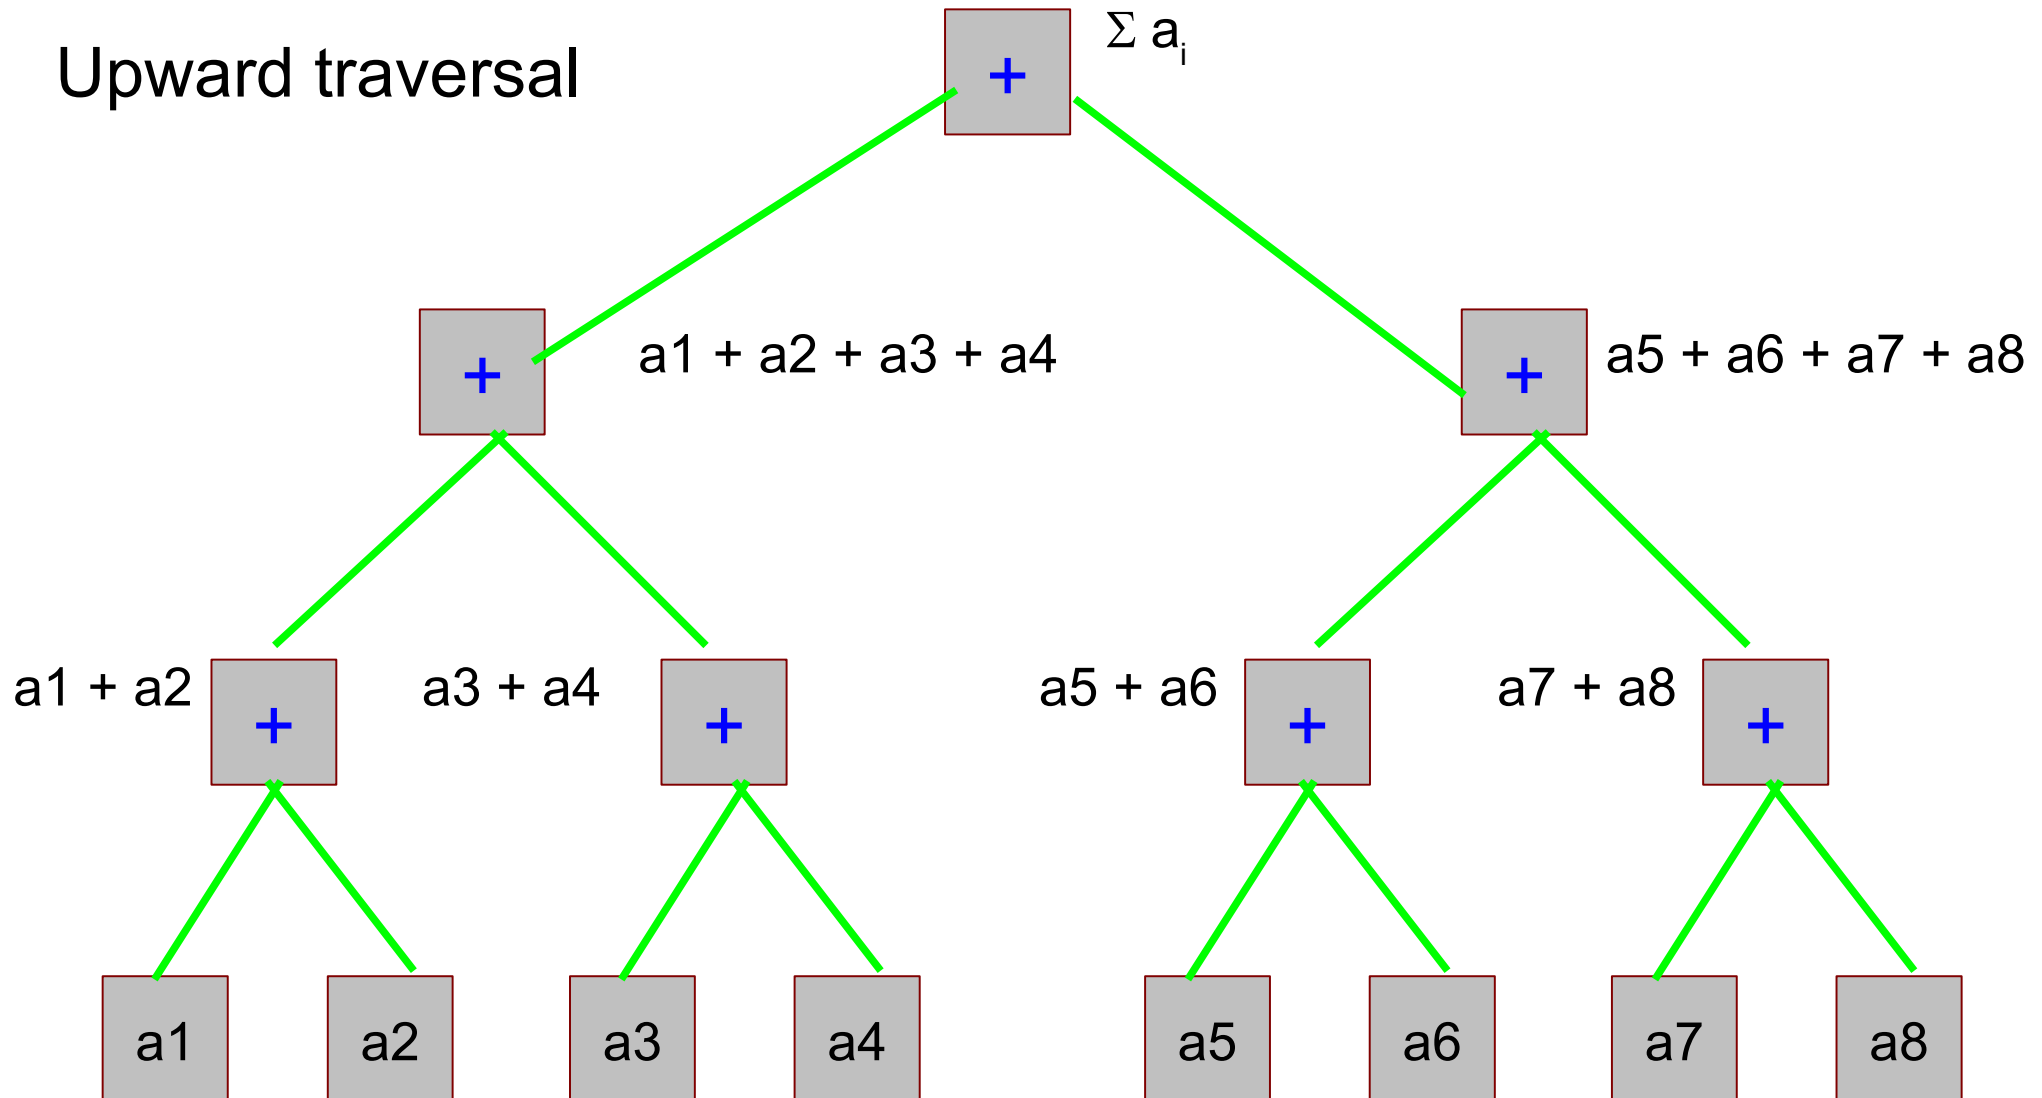


# Balanced Binary Tree – Sum



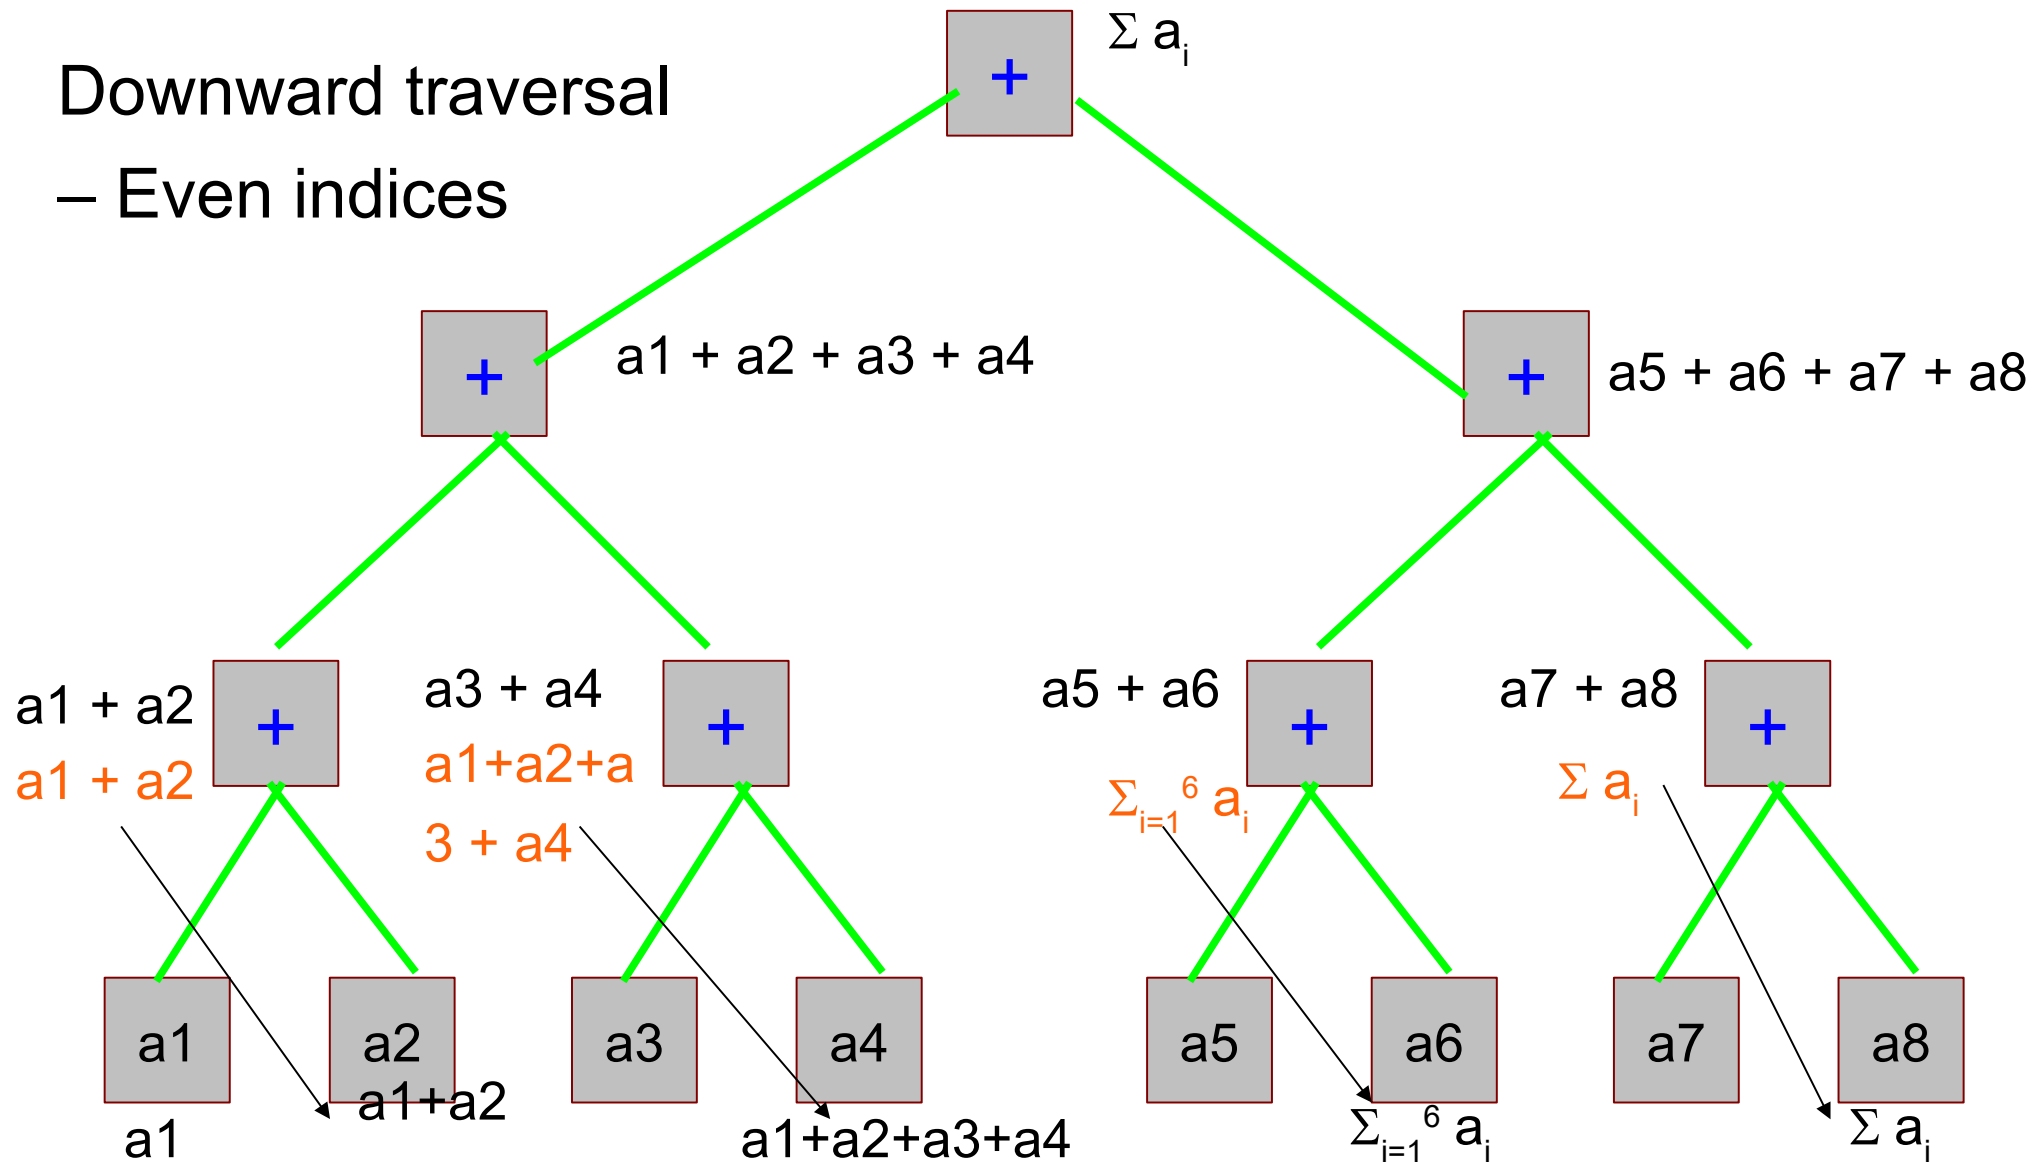
# Balanced Binary Tree – Prefix Sum

Upward traversal




# Balanced Binary Tree – Prefix Sum

Downward traversal  
– Even indices



# Balanced Binary Tree – Prefix Sum

Downward traversal  
– Odd indices



# Balanced Binary Tree – Prefix Sums

- Two traversals of a complete binary tree.
- The tree is only a visual aid.
  - Map processors to locations in the tree
  - Perform equivalent computations.
  - Algorithm designed in the PRAM model.
  - Works in logarithmic time, and optimal number of operations.

//upward traversal

1. for  $i = 1$  to  $n/2$  do in parallel

$$b_i = a_{2i-2} \circ a_{2i}$$

2. Recursively compute the prefix sums of  $B = (b_1, b_2, \dots, b_{n/2})$  and store them in  $C = (c_1, c_2, \dots, c_{n/2})$

//downward traversal

3. for  $i = 1$  to  $n$  do in parallel

$$i \text{ is even} : s_i = c_i$$

$$i = 1 : s_1 = c_1$$

$$i \text{ is odd} : s_i = c_{(i-1)/2} \circ a_i$$

# Analysis of Parallel Algorithms

---

- To analyze parallel algorithms, we rely on asymptotics and recurrences.
- Each operation costs 1 unit, only sequential time needs to be counted. We assume **as many processors as can be used** are available.
- In the prefix sum example, let  $T(n)$  be the time in parallel for an input of size  $n$ .
  - Step 1 can use  $n/2$  processors in parallel each taking 1 unit of time.
  - Step 2 is a recursive call and takes  $T(n/2)$  time.
  - Step 3 uses  $n$  processors each taking 1 unit of time.

# Analysis of Parallel Algorithms

---

- The recurrence relation is:
  - $T(n) = T(n/2) + O(1)$
  - Can ignore effects due to constant factors, such as the difference in the number of processors between steps 1 and 3.
- The solution to the above recurrence is  $T(n) = O(\log n)$ .
- Another parameter of interest in parallel algorithms is the work done.
- Can be stated as the sum of the works done by each of the processors.

# Analysis of Parallel Algorithms

---

- The work done by the prefix algorithm can be expressed by the recurrence
  - $W(n) = W(n/2) + O(n)$ .
  - The  $O(n)$  accounts for the work in the first and the third steps.
  - Solution:  $W(n) = O(n)$ .
- Work done can indicate if the algorithm is doing about the same amount of operations as the best known sequential algorithm.
- Such a parallel algorithm is called an **optimal algorithm**.




# Other Design Paradigms

---

- Partitioning
  - Similar to divide and conquer
  - But **no need** to combine solutions
  - Can treat problems independently and solve in parallel.
  - Example: Parallel merging, searching.

# Merging in Parallel by Partitioning

---



- Two sorted arrays A and B to be merged into C.
- Claim:  $\text{Rank}(x, C) = \text{Rank}(x, A) + \text{Rank}(x, B)$
- For  $x$  in A,  $\text{Rank}(x, A)$  is immediately available. To find  $\text{Rank}(x, B)$  can use binary search in parallel.

# Quick Example

A = [8 10 12 24 ]

B = [15 17 27 32]

Element	8	10	12	24	15	17	27	32
Rank in A	0	1	2	3	3	3	4	4
Rank in B	0	0	0	2	0	1	2	3
Rank in C	0	1	2	5	3	4	6	7

C = [ 8 10 12 15 17 24 27 32 ]

# Merging in Parallel by Partitioning

---

- Time for each binary search is  $O(\log n)$
- Total time for merging =  $O(\log n)$ , the total work is  $O(n \log n)$ .
  - Non optimal as compared to sequential time complexity of  $O(n)$ .
- Can reduce the total work to  $O(n)$ .
  - Induce partitions in the arrays of equal size
  - Rank one element from each partition
  - Use these ranks to find the ranks of the other elements, sequentially.

# An Improved Optimal Algorithm

---

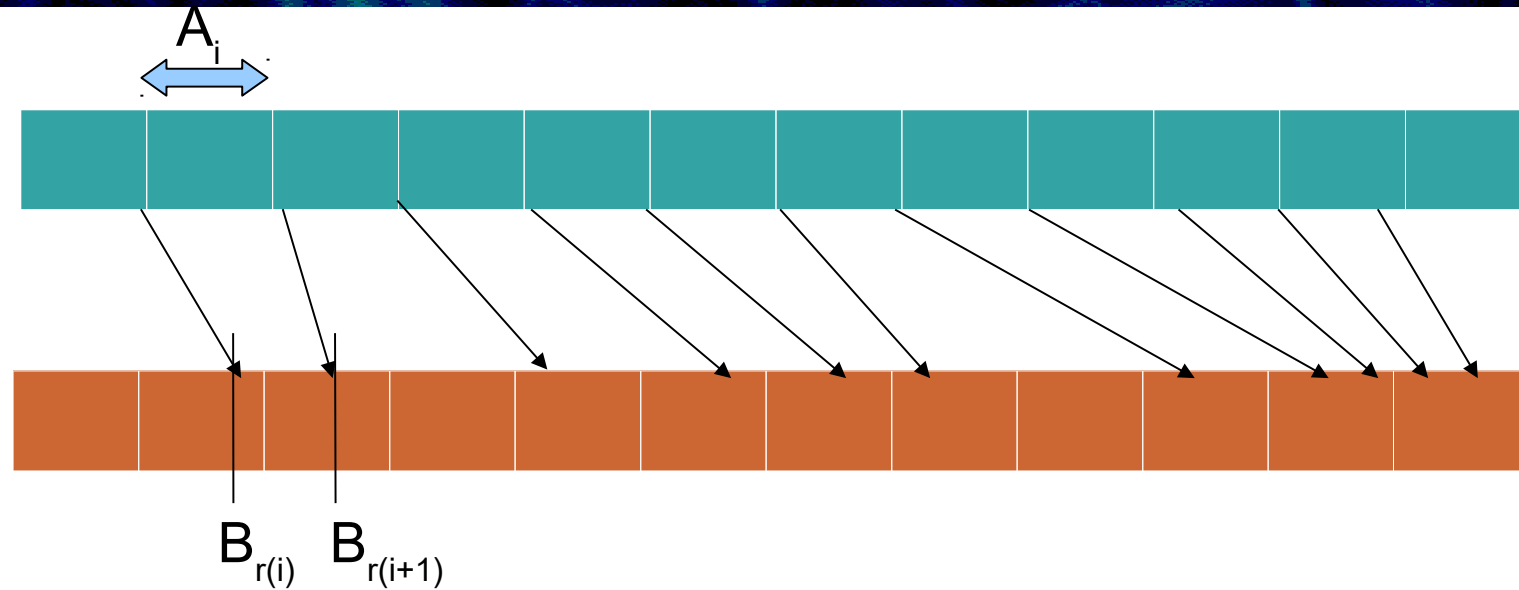
- General technique
  - Solve a smaller problem in parallel
  - Extend the solution to the entire problem.
- For the first step, the problem size to be solved is guided by the factor of non-optimality factor of an existing parallel algorithm.

# An Improved Parallel Algorithm

---

- Our simple parallel algorithm is away from optimality by a factor of  $O(\log n)$ .
- So, we should solve a problem of size  $O(n/\log n)$ .
- For this purpose, we pick every  $\log n^{\text{th}}$  element of  $A$ , and similarly in  $B$ .
- Use the simple parallel algorithm on these elements of  $A$  and  $B$ .
  - Binary search however in the entire  $A$  and  $B$ .

# An Improved Parallel Algorithm



- Let  $A_1, A_2, \dots, A_{n/\log n}$  be the elements of  $A$  ranked in  $B$ .
- These ranks induce partitions in  $B$ .
  - Define  $[B_{r(i)} \dots B_{r(i+1)}]$  as the portion of  $B$  so that  $[A(i) \dots A(i+1)]$  have ranks in.
- Can therefore merge  $[A(i) \dots A(i+1)]$  with  $[B_{r(i)} \dots B_{r(i+1)}]$  sequentially.

# An Improved Parallel Algorithm


---

- Such sequential merges can happen in parallel, at each index of  $A[i]$ .
- Time taken for the sequential merge is  $O(\log n + B_{r(i+1)} - B_{r(i)})$ .
- Time:
  - Binary search:  $O(\log n)$ , with  $n/\log n$  processors.
  - Sequential merge:  $O(\log n)$ , subject to certain conditions. There are also  $n/\log n$  such merges in parallel.
- Work:
  - There are  $n/\log n$  binary searches in parallel. Work =  $O(n)$ .
  - For the sequential merges too, work =  $O(n)$ .




# An Improved Parallel Algorithm

---



- What if  $[B_{r(i)} \dots B_{r(i+1)}]$  has a size of more than  $\log n$ ?
- The situation can be addressed
  - Pick equally spaced, no more than  $\log n$ , spaced items in  $[B_{r(i)} \dots B_{r(i+1)}]$ .
  - Rank these in  $[A_i \dots A_{i+1}]$ .

# An Improved Parallel Algorithm



- What if  $[B_{r(i)} \dots B_{r(i+1)}]$  has a size of more than  $\log n$ ?
- The situation can be addressed
  - Pick equally spaced, no more than  $\log n$ , spaced items in  $[B_{r(i)} \dots B_{r(i+1)}]$ .
  - Rank these in  $[A_i \dots A_{i+1}]$ .

# Final Result

---

- Can merge two sorted arrays of size  $n$  in time  $O(\log n)$  with work  $O(n)$ .
  - Need CREW model, for binary searches.
- Can improve further, we will see later.
- The technique to achieve optimality is a general technique, with several applications. We will see more applications of this later.

# A Further Improvement

---

- Where is the scope for improvement?
- Each binary search takes  $O(\log n)$  time, and we also have  $O(n/\log n)$  subproblems each of size  $O(\log n)$ .
- To get further improvements, we should look at both aspects.
- Can we search faster? Parallel?

# A Further Improvement

---

- Parallel search first.
- Consider a sorted array  $A$  of  $n$  element and we want to search for an element  $x$ .
- Given  $p$  processors, we can always search at positions (indices)  $1, n/p, 2n/p, \dots, n$ .
- Record the result of each comparison as a 1 or 0 with 1 for position  $i$  indicating that  $A[i] < x$  and 0 indicating that  $A[i] \geq x$ .
- The sequence of  $p$  results will have :
  - Either all 1's
  - Either all 0's
  - A shift from 1's to 0's

# A Further Improvement

---

- Parallel search first.
- Consider a sorted array  $A$  of  $n$  element and we want to search for an element  $x$ .
- Given  $p$  processors, we can always search at positions (indices)  $1, n/p, 2n/p, \dots, n$ .
- Record the result of each comparison as a 1 or 0 with 1 for position  $i$  indicating that  $A[i] < x$  and 0 indicating that  $A[i] \geq x$ .
- The sequence of  $p$  results will have :
  - Either all 1's :  $x$  is not in  $A$
  - Either all 0's :  $x$  is not in  $A$
  - A shift from 1's to 0's :  $x$  is likely in the  $n/p$  segment corresponding to the shift from 1 to 0.

# Search in Parallel

---

- We can identify the next step depending on the three cases.
  - Either all 1's :  $x$  is not in  $A$
  - Either all 0's :  $x$  is not in  $A$
  - A shift from 1's to 0's :  $x$  is likely in the  $n/p$  segment corresponding to the shift from 1 to 0.
    - Therefore, search recursively in the corresponding segment of size  $n/p$  while still using  $p$  processors.
- The recurrence relation for the time taken is
  - $T(n) = T(n/p) + O(1)$ , for a solution of  $T(n) = O(\log_p n)$ .
- The work done has the recurrence  $W(n) = p \cdot W(n/p) + O(p)$ , for a solution of  $W(n) = O(n/p)$ .

# Search in Parallel

---

- Consider typical values of  $p$ .
- For  $p = O(1)$ , no change in time taken asymptotically.
- For  $p = O(\log n)$ , the time taken is  $O(\log n / \log \log n)$ .
- For  $p = O(n^{1/2})$ , the time taken is  $O(\log n / \log n^{1/2}) = O(1)$ !
  - Of course, looks like wasteful from a work point of view.
  - Let us see what it is good for!



# From Parallel Search to Merge

---

- Recall our idea to arrive at an optimal algorithm to merge two sorted arrays  $A$  and  $B$ .
- We rank a few elements of  $A$  in  $B$  to partition  $B$  into sub-arrays.
- Let us consider ranking  $n^{1/2}$  elements of  $A$  in  $B$ .
- We have  $n$  processors, so each search can use  $n^{1/2}$  processors!
- Each search now finishes in  $O(1)$  time.

# From Parallel Search to Merge

---

- Let us consider ranking  $n^{1/2}$  elements of A in B.
- We have  $n$  processors, so each search can use  $n^{1/2}$  processors!
- Each search now finishes in  $O(1)$  time.
- There is a downside however.
- The partitions of A are now much longer at  $n^{1/2}$  each.
- The partitions of B are like in the earlier case, unknown.

# From Parallel Search to Merge

---

- There is a downside however.
- The partitions of A are now much longer at  $n^{1/2}$  each.
- The partitions of B are like in the earlier case, unknown.
- But, can use recursion to make further progress.

# From Parallel Search to Merge

---

- The partitions of A are now much longer at  $n^{1/2}$  each.
- The partitions of B are like in the earlier case, unknown.
- But, can use recursion to make further progress.
- Recursively apply the same procedure on each partition of A into the corresponding partition of B.
- Notice that each part of A is only  $n^{1/2}$  in size.
- We want to rank  $n^{1/4}$  element of each part of A into the corresponding B.

# From Parallel Search to Merge

---

- The recurrence relation guiding this process is captured by  $T(n, m) = \max_i T(n^{1/2}, m_i) + O(1)$ .
  - In the above,  $n$  and  $m$  refer to the length of  $A$  and  $B$  respectively.
  - And,  $m_i$  refers to the length of the  $i^{\text{th}}$  partition of  $B$ .
- Can show that  $T(n, m) = O(\log \log n)$ .
- Once recursion ends, each partition of  $A$  and partitions of  $B$  will be  $O(\log \log n)$  long, and we merge them sequentially.

# The Power of CRCW – Minima

---

- Two points of interest
  - Illustrate the power of CRCW models
  - Illustrate another optimality technique.
- Find minima of  $n$  elements.
  - Input: An array  $A$  of  $n$  elements
  - Output: The minimum element in  $A$ .
- From what we already know:
  - Standard sequential algorithm not good enough
  - Can use an upward traversal, with  $\min$  as the operator at each internal node. Time =  $O(\log n)$ , work =  $O(n)$ .

# The Power of CRCW – Minima

---

- Our solution steps:
  - Design a  $O(n^2)$  work,  $O(1)$  time algorithm.
  - Gain optimality by sacrificing runtime to  $O(\log \log n)$ .

# An $O(1)$ Time Algorithm

---

	12	17	8	18	26
12	--	1	0	1	1
17	0	--	0	1	1
8	1	1	--	1	1
18	0	0	0	--	1
26	0	0	0	0	--

- Use  $n^2$  processors.
- Compare  $A[i]$  with  $A[j]$  for each  $i$  and  $j$ .
- Now can identify the minimum.



# An $O(1)$ Time Algorithm

---

	12	17	8	18	26
12	--	1	0	1	1
17	0	--	0	1	1
8	1	1	--	1	1
18	0	0	0	--	1
26	0	0	0	0	--

- Use  $n^2$  processors.
- Compare  $A[i]$  with  $A[j]$  for each  $i$  and  $j$ .
- Now can identify the minimum.
  - How?

# An $O(1)$ Time Algorithm

---

	12	17	8	18	26
12	--	1	0	1	1
17	0	--	0	1	1
8	1	1	--	1	1
18	0	0	0	--	1
26	0	0	0	0	--


- Use  $n^2$  processors.
- Compare  $A[i]$  with  $A[j]$  for each  $i$  and  $j$ .
- Now can identify the minimum.
  - How?
- Where did we need the CRCW model?

# Towards Optimality

---

- The earlier algorithm is heavy on work.
- To reduce the work, we proceed as follows.
- We derive an  $O(n \log \log n)$  work algorithm running in  $O(\log \log n)$  time.
- For this, use a doubly logarithmic tree.
  - Defined in the following.

# Doubly Logarithmic Tree



- Let there be  $n = 2^{2^k}$  leaves, the root is level 0. The root has  $\sqrt{n} = 2^{2^{k-1}}$  children.
- In general, a node at level  $i$  has  $2^{2^{k-i-1}}$  children, for  $0 \leq i \leq k-1$ .
- Each node at level  $k$  has two leaf nodes as children.

# Doubly Logarithmic Tree

---

- Some claims:
  - Number of nodes at level  $i$  is  $2^{2^k} - 2^{2^k-i}$ .
  - Number of nodes at the  $k$ th level is  $n/2$ .
  - Depth of a doubly logarithmic tree of  $n$  nodes is  $k+1 = \log \log n + 1$ .
- To compute the minimum using a doubly logarithmic tree:
  - Each internal node performs the min operation does not suffice.
  - Why?

# Minima Using the Doubly Logarithmic Tree

---

- Intuition:
  - Should spend only  $O(1)$  time at each internal node.
  - Use the  $O(1)$  time algorithm at each internal node.
- At each internal node of level  $i$ , if there are  $c_i$  children, use  $c_i^2$  processors.
  - Minima takes  $O(1)$  time at each level.
  - Also, No. of nodes at level  $i$  x No. of processors used =  $2^{2k} - 2^{k-i} \cdot (2^{2^{k-i-1}})^2 = 2^{2k} = n$ .

# Minima Using a Doubly Logarithmic Tree

---

- Second, slightly improved result:
  - With  $n$  processors, can find the minima of  $n$  numbers in  $O(\log \log n)$  time.
  - Total work =  $O(n \log \log n)$ .
- Still suboptimal by a factor of  $O(\log \log n)$ .
- We now introduce a technique to achieve optimality.

# Accelerated Cascading

---

- Our two algorithms:
  - Algorithm 1: A slow but optimal algorithm.
    - Binary tree based:  $O(\log n)$  time,  $O(n)$  work.
  - Algorithm 2: A fast but non-optimal algorithm
    - Doubly Logarithmic tree based:  $O(\log \log n)$  time,  $O(n \log \log n)$  work.
- The **accelerated cascading** technique suggests combining two such algorithms to arrive at an optimal algorithm
  - Start with the slow but optimal algorithm till the problem is small enough
  - Switch over to the fast but non-optimal algorithm.



# Accelerated Cascading

---

- The binary tree based algorithm starts with an input of size  $n$ .
- Each level up the tree reduces the size of the input by a factor of 2.
- In  $\log \log \log n$  levels, the size of the input reduces to  $n/2^{\log \log \log n} = n/\log \log n$ .
- Now switch over to the fast algorithm with  $n/\log \log n$  processors, needing  $O(\log \log (n/\log \log n))$  time.

# Final Result

---

- Total time =  $O(\log \log \log n) + O(\log \log n)$ .
- Total work =  $O(n)$ .
- Need CRCW model.
- Where did we need the CRCW model?