

CSE 475: Statistical Methods in AI

Monsoon 2019

SMAI-M-2019 11: Linear Perceptrons

Lecturer: C. V. Jawahar

Date: 9 Sep 2019

11.62 GD for Classification

We know the gradient descent equation as

$$\mathbf{w}^{k+1} \leftarrow \mathbf{w}^k - \eta \nabla J \quad (11.16)$$

We start with an initialization \mathbf{w}^0 and update the weights/vector until we get the separating hyperplane (or until it converges).

Let us specially focus on the classification problem today. We wish to use a loss that measures classification accuracy. That is, the percentage of samples misclassified.

Let us now consider the objective as

$$J = \frac{1}{N} \sum_{i=1}^N (1 - y_i \cdot f(\mathbf{w}, \mathbf{x}_i)) \quad (11.17)$$

where $y_i \in \{-1, +1\}$, and $f(\mathbf{x}_i) = \text{sign}(\mathbf{w}^T \mathbf{x}_i)$. This provides a “unit” loss for all the misclassifications and “zero” loss for all the correct classifications. (see notes elsewhere for details.) We are not using $\text{sign}()$. Our ideal goal could have been to minimize this loss.

- Q: Do we have to divide the loss by 2 to get unit loss per sample?
- Q: What makes this problem non-differentiable?

However we have a serious problem. This is not differentiable. Then, what do we do? We can address this in two different ways. This leads to some popular algorithms:

- Perceptron Algorithm
- Logistic Regression
- Support Vector Machines (SVM)

11.63 Perceptron Algorithm

Perceptron algorithm makes the assumption that all the samples are linearly separable. In other words $\exists \mathbf{w}$ such that $y_i \cdot f(\mathbf{w}, \mathbf{x}_i) = +1 \forall i$.

- Q: Can every set of (i) Two (ii) Three (iii) Four points be linearly separable in 2D?
- Q: Can any specific set of (i) Two (ii) Three (iii) Four be linearly separable in 2D for all potential labelings of these samples?
- Q: Can every set of (i) Two (ii) Three (iii) Four points be linearly separable in 2D for all potential labelings of these points?
- Note: You may know about the Ex-OR case. You may also want to know about VC-Dimension to appreciate the above questions? Why are they important in ML?

Let us now look at an alternate loss function. We define the loss/objective as

$$J = \sum_{\mathbf{x}_i \in \mathcal{E}} -y_i \cdot \mathbf{w}^T \mathbf{x}_i \quad (11.18)$$

where \mathcal{E} is the set of misclassified samples. i.e.,

$$\mathcal{E} = \{\mathbf{x}_i | \mathbf{w}^T \mathbf{x}_i < 0\}$$

Note that J is always non-negative. It can become zero when all samples are correctly classified or \mathcal{E} is empty.

How is this loss function different?

- The summation is only over the misclassified samples. This does not change anything really. (We had zero loss for all the correctly classified ones anyway).
- We then have a “non-unity” loss for all the misclassified ones. (This is different from the previous loss). It is proportional to how far it is from the line/plane/hyperplane.
- Q: Why not then we use the following objective? (Demonstrate with an example)

$$J = \frac{1}{N} \sum_{i=1}^N -y_i \cdot \mathbf{w}^T \mathbf{x}_i$$

- In our objective, the farther the sample, the more the loss is. (Indeed, not very ideal!!). *Q: Why do you think this is not ideal? Is it really non ideal if our objective sums only over \underline{E}*

Pleasantly, this is now differentiable. That is an advantage. Also when all samples are correctly classified (when the problem is linearly separable), the loss becomes zero (\mathcal{E} becomes empty set.) and our algorithm will converge (loss is zero, derivative is also zero.)

There is also an advantage with this loss. A sample that is far from the line/plane will pull/push/rotate the line/plane more than one that is very near the line. This will help in faster convergence.

Let us now re-write our gradient descent equation as:

$$\mathbf{w}^{k+1} \leftarrow \mathbf{w}^k + \eta \sum_{\mathbf{x}_i \in \mathcal{E}} y_i \mathbf{x}_i \quad (11.19)$$

Q: Verify this.

We start with a single sample version of the perceptron. i.e., the perceptron algorithm is now:

- Initialize $k = 0$, \mathbf{w}^0
- While \mathcal{E} is not empty
 - Pick an arbitrary element \mathbf{x}_j from \mathcal{E}
 - $\mathbf{w}^{k+1} \leftarrow \mathbf{w}^k + y_j \mathbf{x}_j$
 - $k \leftarrow k + 1$
- Return \mathbf{w}_k

Note that (i) We assumed the learning rate $\eta = 1$ (ii) We update the \mathbf{w} for each sample, than for a batch.

For a batch of samples, we can define the weight update as:

$$\mathbf{w}^{k+1} \leftarrow \mathbf{w}^k + \eta \sum_{\mathbf{x}_j \in \mathcal{E}} y_j \mathbf{x}_j$$

Another way we could define the weight update is with the help of desired/target (t) and output (o). In this situation, the update rule is:

$$\mathbf{w}^{k+1} \leftarrow \mathbf{w}^k + \eta \sum_{i=1}^N (t_i - o_i) \mathbf{x}_i$$

- Q: prove that both the update rules are the same (with a scale change in η).

Note that when desired and predicted outputs are same, $t - o$ is zero. Else it is either Positive or negative (when y_i as well as $t - o$ will have the same sign when the sample is misclassified.)

$$\mathbf{w}^{k+1} \leftarrow \mathbf{w}^k + \eta \sum_{i=1}^N (t_i - o_i) \mathbf{x}_i \quad (11.20)$$

Note that η is a learning rate and it could absorb any scaling. (η in all these equations need not be identical, see yourself, if they differ by a scale factor.). Here the summation is over all the samples. Note that, this does not change the update rule. The additional terms are zero.

Algorithm now has the following steps:

1. Initialize \mathbf{w} , $k=0$

2. We update the \mathbf{w} as

$$\mathbf{w}^{k+1} \leftarrow \mathbf{w}^k + \eta \sum_{i=1}^N (t_i - o_i) \mathbf{x}_i$$

or

$$\mathbf{w}^{k+1} \leftarrow \mathbf{w}^k + \eta \sum_{\mathbf{x}_i \in \mathcal{E}} y_i \mathbf{x}_i$$

3. $k \leftarrow k + 1$

4. Repeat steps 2-4

- until \mathcal{E} is empty. (the ideal termination criteria for perceptron algorithm) Or
- until the change in weight is small (say less than θ).

11.64 More Details

11.64.1 Delta Rule

Let us consider a “regression” problem with o_i as $\mathbf{w}^T \mathbf{x}_i$. Consider a least square objective as

$$J = \sum_{i=1}^N (t_i - o_i)^2$$

The gradient descent update rule is:

$$\mathbf{w}^{k+1} \leftarrow \mathbf{w}^k + \eta \sum_{i=1}^N (t_i - o_i) \mathbf{x}_i$$

- Q: is it really the same as we saw early? Not really!!

11.64.2 Back to Perceptrons

Is it really possible to arrive the gradient descent update rule as Equation 11.20 from the objective like:

$$J = \sum_{i=1}^N (t_i - o_i)^2$$

with $o_i = \text{sign}(\mathbf{w}^T \mathbf{x}_i)$? There is a small catch.

With small changes in the \mathbf{w} or the line, the J is not changing. This is not very appropriate to look at this as gradient descent objective. Let us modify this as

$$J = \sum_{i=1}^N (t_i - o_i)^2 (-\mathbf{w}^T \mathbf{x}_i)$$

This additional terms pulls (or pushes) *proportionally*. Let us now rewrite this J as sum of two parts one over \mathcal{E} and the other on not in \mathcal{E} .

$$J = J_1 + J_2 = \sum_{\mathbf{x}_i \in \mathcal{E}} (t_i - o_i)^2 (-\mathbf{w}^T \mathbf{x}_i) + \sum_{\mathbf{x}_i \text{ not in } \mathcal{E}} (t_i - o_i)^2 (-\mathbf{w}^T \mathbf{x}_i)$$

$$J = J_1 + J_2 = \sum_{\mathbf{x}_i \in \mathcal{E}} (2 \cdot t_i)^2 (-\mathbf{w}^T \mathbf{x}_i) + \sum_{\mathbf{x}_i \text{ not in } \mathcal{E}} 0 \times (\mathbf{w}^T \mathbf{x}_i)$$

We know that J_2 is zero. When $\mathbf{x}_i \in \mathcal{E}$, $(t_i - o_i)$ is $2t_i$ i.e., $2y_i$.

$$\nabla J = 2 \sum_{i=1}^N (t_i - o_i) \mathbf{x}_i$$

Note that:

- Either $(t_i - o_i)$ is zero
- Or we know that $(t_i - o_i)$ is $2t_i$ and $\frac{\partial -\mathbf{w}^T \mathbf{x}_i}{\partial \mathbf{w}} = -\mathbf{x}_i$.

This leads to:

$$\mathbf{w}^{k+1} \leftarrow \mathbf{w}^k - \eta \nabla J$$

$$\mathbf{w}^{k+1} \leftarrow \mathbf{w}^k + \eta \sum_{i=1}^N (t_i - o_i) (-\mathbf{x}_i)$$

$$\mathbf{w}^{k+1} \leftarrow \mathbf{w}^k + \eta \sum_{\mathbf{x}_i \in \mathcal{E}} 2 \cdot t_i \mathbf{x}_i$$

$$\mathbf{w}^{k+1} \leftarrow \mathbf{w}^k + \eta \sum_{\mathbf{x}_i \in \mathcal{E}} y_i \mathbf{x}_i$$

(with changes in scale for learning rate η)

11.65 Discussions and Examples

Let us consider some simple situations first and see how the perceptron algorithm behaves.

Let us simplify the update rule first. Assume $\eta = 1$. Let us assume that there is only one sample. Also let us assume that \mathbf{x} has only one dimension i.e., scalar. This simplifies to:

$$w^{k+1} \leftarrow w^k + (t_i - o_i) x_i$$

Now let us consider some situations.

- When there is no misclassification. i.e., $t_i = o_i$. In this case, w does not change.
- Let us consider $t_i = +1$, $o_i = -1$ and x_i is positive. Sample is misclassified. Which means $w^k \cdot x_i$ is negative. (x_i is fixed and positive.) This means w^k is negative. We need to increase w^k to (positive to) minimize/avoid misclassification. If we substitute the values/signs in the above equation, we see that perceptron algorithm does this exactly.
- Now Let us consider $t_i = +1$, $o_i = -1$ and x_i is negative. Sample is misclassified. Which means $w^k \cdot x_i$ is negative. (x_i is fixed and negative.) Therefore, w^k is positive. We need to decrease w^k to minimize/avoid misclassification. If we substitute the values/signs in the above equations, we see that perceptron algorithm does this exactly.
- Q: Convince yourself for all other cases also.

11.65.1 Example/Problem in 2D

Q: Consider a set of vectors in 2D.

$$\{[1, 1]^T, [1, 3]^T, [2, 1]^T, [2, 2]^T,$$

$$[-1, -1]^T, [-1, -3]^T, [-2, -1]^T, [-2, -2]^T\}$$

The first four are from class 1 and the rest four are from class 2.

- Plot the samples in a 2D plane with “x” for positive classes and “o” for negative class. Is this set linearly separable?
- Start with a random vector for \mathbf{w} and show that the perceptron algorithm converges to a separating line/plane for different values of η . (make sure that you start with a vector that has error!!).
- Why is that the final answers are different for different initializations/ η ? Then which is the best solution?

11.66 Some Theoretical Results

Perceptron algorithm has a number of interesting theoretical properties/results. A summary is below.

- If there exist a set of weights that are consistent with the data, the perceptron algorithm will converge.
- If the training data is not Linearly Separable, the perceptron algorithm will eventually repeat the same set of weights and thereby enter an infinite loop.

- If the training data is linearly separable, algorithm will converge in a maximum of M steps. (See more below for the bounds).
- Every boolean function can be represented by some network of perceptrons only two levels deep.

11.67 Closer Look at the Theoretical Results

Problem Our problem is to design a classification algorithm for N samples $\mathcal{D} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$ where $\mathbf{x}_i \in R^d$ and $y_i \in \{-1, +1\}$

We make the assumption that the data is linearly separable. What does it mean?

This simply means that there exists a solution $\mathbf{w}^* \in R^d$ such that $\|\mathbf{w}^*\| = 1$ and a small quantity γ such that

$$y_i(\mathbf{w}^{*T} \mathbf{x}_i) > \gamma$$

for all $i = 1, \dots, N$. Note that when the $\text{sign}(\mathbf{w}^T \mathbf{x})$ is positive and y_i is $+1$, the product is positive. Similarly when the sign $(\mathbf{w}^T \mathbf{x})$ is negative and y_i is -1 , the product is still positive. In other-words, when a sample is classified correctly, the product is always positive and the inequality above is true for a small positive quantity γ .

We also make an assumption that samples are bounded. Not a strong assumption. What does it mean?

$$\|\mathbf{x}_i\| \leq R \quad \forall i$$

11.67.1 Summary of Perceptron Algorithm Assumptions

We start with a random initialization and update solution based on the misclassified samples. We repeat this until we find that all the samples are correctly classified.

We make two assumptions.

1. $y_i(\mathbf{w}^{*T} \mathbf{x}_i) > \gamma$ for all i
2. $\|\mathbf{x}_i\| \leq R$ for all i .
3. $\|\mathbf{w}^*\| = 1$

Where \mathbf{w}^* is the final/optimal vector. γ is a positive quantity that ensure separability and also a margin. The first one says that the samples are separable. The second one says that all the samples are bounded. (not of infinite magnitude.). Note that γ and R are positive.

11.68 Convergence of Perceptron

Now our main result is *The Perceptron Learning Algorithm makes at most $\frac{R^2}{\gamma^2}$ iterations and converge. It returns a separating hyperplane after that.*

We know the perceptron update equation as

$$\mathbf{w}^{k+1} = \mathbf{w}^k + y_i \mathbf{x}_i$$

This assumes $\eta = 1$. Also this is a single sample update rule.

Let us assume that we start with \mathbf{w}^0 as a zero vector $\mathbf{0}$.

Let us multiply \mathbf{w}^* on both sides.

$$(\mathbf{w}^{k+1})^T \mathbf{w}^* = (\mathbf{w}^k)^T \mathbf{w}^* + y_i (\mathbf{x}_i)^T \mathbf{w}^*$$

since the last term is greater than γ (first assumption).

$$(\mathbf{w}^{k+1})^T \mathbf{w}^* > (\mathbf{w}^k)^T \mathbf{w}^* + \gamma$$

From induction

$$(\mathbf{w}^{k+1})^T \mathbf{w}^* > k\gamma$$

We know that $\mathbf{a}^T \mathbf{b} < \|\mathbf{a}\| \cdot \|\mathbf{b}\|$. Therefore,

$$(\mathbf{w}^{k+1})^T \mathbf{w}^* < \|\mathbf{w}^*\| \cdot \|\mathbf{w}^{k+1}\| = \|\mathbf{w}^{k+1}\|$$

or finally

$$\|\mathbf{w}^{k+1}\| > k\gamma \quad \text{and} \quad \|\mathbf{w}^{k+1}\|^2 > k^2 \gamma^2 \quad (11.21)$$

We have now a lower bound on $\|\mathbf{w}^{k+1}\|^2$

Let us also try to get an upper bound for $\|\mathbf{w}^{k+1}\|^2$

$$\mathbf{w}^{k+1} = \mathbf{w}^k + y_i \mathbf{x}_i$$

Therefore:

$$\|\mathbf{w}^{k+1}\|^2 = \|\mathbf{w}^k + y_i \mathbf{x}_i\|^2$$

$$\|\mathbf{w}^{k+1}\|^2 = \|\mathbf{w}^k\|^2 + \|y_i \mathbf{x}_i\|^2 + 2\|\mathbf{w}^k\| \mathbf{x}_i^T y_i$$

Note that the last term is negative. Since the updates happen only when there is an error in the sample/classification.

$$\|\mathbf{w}^{k+1}\|^2 \leq \|\mathbf{w}^k\|^2 + \|\mathbf{x}_i\|^2$$

$$\|\mathbf{w}^{k+1}\|^2 \leq \|\mathbf{w}^k\|^2 + R^2$$

$$\|\mathbf{w}^{k+1}\|^2 \leq kR^2$$

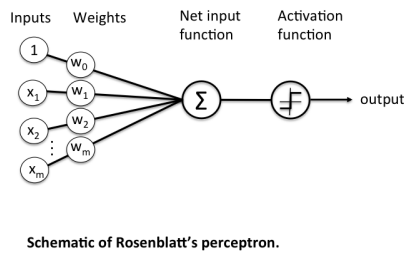


Figure 11.6: A pictorial representation of the neuron. Often a single large circle is shown for a neuron wherein the weighted addition and nonlinearity is combined.

per sample and update for each sample. A sample may be seen multiple times (in a cyclic manner).

- Batch version: Compute the gradient for the batch (full set) and update once.
- Mini-Batch: Instead of taking the full batch, take a smaller batch. This is useful when the data is large.
- Stochastic: Samples in online or mini-batch are selected randomly.

Q: Do write the pseudo code for each of these. Make sure your pseudo code is very close to the programming language that you use.

11.68.1 Combining the Bounds

We can now write

$$k^2 \gamma^2 < \|\mathbf{w}^{k+1}\|^2 \leq kR^2$$

Combining the upper and lower bounds.

$$k < \frac{R^2}{\gamma^2}$$

11.69 Neural Networks View Point

A popular view point of the perceptron is to appreciate it as a “neuron”. Though the story has origins in our attempts to understand and reverse engineer human brain and perception skills, it is much easier to appreciate it as a powerful mathematical model at this stage.

A neuron accepts multiple inputs. Weigh each one. Add the inputs. Pass through a nonlinearity. In the case of perceptron, the nonlinearity is a step like nonlinearity. See also figure.

We will revisit the neural network view point at a later stage when we discuss “multi layer perceptrons”.

11.70 Variations in Gradient Descent

We now know: (i) How to define a loss function (ii) how to optimize it with gradient descent update equations.

There are in fact many variations in the implementation.

- Single sample or online version: which assumes that samples come one by one. This computes gradient