

Topics in machine learning: Naresh Manwani

Siddharth Bhat

Monsoon 2019

Contents

1	Policy iteration	5
1.1	Value iteration algorithm	5
2	Monte carlo methods for MDP	7
2.1	Naive	7
2.1.1	First visit monte carlo policy evaluation	7
2.1.2	Every visit monte carlo	8
2.2	Soft policy	9
2.3	On policy first-visit MC control, for ϵ -soft policy	9
2.4	Off policy	10

Chapter 1

Policy iteration

$$\pi_{k+1}(s) = \arg \max_{a \in \mathcal{A}(s)} r(s, a) + \gamma \sum_s P(s'|s, a) v_{\pi_k}(s')$$

Theorem 1 *The policy iteration algorithm generates a sequence of policies with non-decreasing state values. That is, $V^{\pi_{k+1}} \geq V^{\pi_k}$, $V^\pi \in \mathbb{R}^n$, is the vector of state values for state π*

Proof 1 F^{π_k} is the bellman expectation operator (?)

Since V^{π_k} is a fixed point of F^{π_k} ,

$$V^{\pi_k} = F^{\pi_k}(V^{\pi_k}) \leq F(V^{\pi_k}) \quad (\text{upper bounded by max value})$$

$$F(V^{\pi_k}) = F^{\pi_{k+1}}(V^{\pi_k}) \quad (\text{By defn of policy improvement step})$$

$$V^{\pi_k} \leq F^{\pi_{k+1}}(V^{\pi_k}) \quad (\text{eqn 1})$$

$$F^{\pi_{k+1}}(V^{\pi_k}) \leq (F^{\pi_{k+1}})^2(V^{\pi_k}) \quad (\text{Monotonicity of } F^{\pi_{k+1}})$$

$$\forall t \geq 1, F^{\pi_{k+1}}(V^{\pi_k}) \leq (F^{\pi_{k+1}})^t(V^{\pi_k}) \quad (\text{Monotonicity of } F^{\pi_{k+1}})$$

$$F^{\pi_{k+1}}(V^{\pi_k}) \leq (F^{\pi_{k+1}})^t(V^{\pi_k}) \leq V^{\pi_{k+1}} \quad (\text{Contraction mapping, } V^{\pi_{k+1}} \text{ is fixed point})$$

$$V^{\pi_k} = F^{\pi_{k+1}}(V^{\pi_k}) \leq V^{\pi_{k+1}}$$

For a set of actions \mathcal{A} and a set of states \mathcal{S} , the total number of policies is $|\mathcal{A}^{\mathcal{S}}|$. The number of computations per iteration is $O(|\mathcal{S}|^3)$. So the loose upper bound is $O(|\mathcal{S}|^3 \times |\mathcal{A}^{\mathcal{S}}|)$.

1.1 Value iteration algorithm

```
let v n s = max [r s a + gamma * sum [(p s' s a) * v (n-1) s' | s' <- ss] | a <- as]
let vs = [v i | i <- [0..]]
-- / L infinity
let norm v v' = max [(v s - v' s) | s <- ss]
let out = head $
  dropWhile (\v v' -> norm (v' - v) < eps * (1 - gamma) / (2 * gamma)) $
  zip vs (tail vs)
let policy s = argmax as $ \a ->
  r s a + gamma * sum [(p s' s a) * out s' | s' <- ss]
```

Theorem 2 For the series V_n and the policy π_ϵ computed by the value iteration algorithm, then:

$$\forall \epsilon > 0, \exists n_0 \in \mathbb{N}, \forall n \geq n_0, \|V_{n+1} - V_n\|_\infty \leq \frac{\epsilon(1-\gamma)}{2\gamma}$$

Proof 2 We need to show that the sequence $\{V_n\}_{n=0}^\infty$ is a Cauchy sequence. This has been proven before by the use of contraction mapping. Thus, for a given $\epsilon' \geq 0, \exists n_0 \in \mathbb{N}, \forall n \geq n_0, \|V_{n+1} - V_n\|_\infty \leq \epsilon'$ by Cauchy sequence. So, pick $\epsilon' = \frac{\epsilon(1-\gamma)}{2\gamma}$, and the proof immediately follows.

Theorem 3 If $\|V_{n+1} - V_n\|_\infty \leq \frac{\epsilon(1-\gamma)}{2\gamma}$, then $\|V_{n+1} - V^*\|_\infty < \epsilon/2$

Proof 3

$$\begin{aligned} \|V_{n+1} - V^*\| &= \|V_{n+1} - FV_{n+1} + FV_{n+1} - V^*\| \leq \|V_{n+1} - FV^*\| + \|FV_n - V_n\| && (\text{triangle inequality}) \\ &\leq \|V_{n+1} - FV^*\| + \gamma\|V_{n+1} - V^*\| \\ &\leq \gamma\|V_{n+1} - V_n\| + \gamma\|V_{n+1} - V^*\| \\ (1-\gamma)\|V_{n+1} - V^*\| &\leq \gamma\|V_n - V_{n+1}\| && (\text{how?}) \\ \implies \dots \end{aligned}$$

It appears that V^{π_ϵ} is just V_{n+2} ??

Theorem 4 The policy π_ϵ is ϵ -optimal: $\|V^* - V^{\pi_\epsilon}\| \leq \epsilon$

Chapter 2

Monte carlo methods for MDP

For dynamic programming, we needed to know the transition probability distribution $P(s, a, s')$, nor the reward function $r(s, a)$.

In the monte carlo methods, we assume that we do not know the transition probability distribution. We rely only on simulations.

This samples over *episodes* for a fixed policy: sequences of states, actions, and rewards.

2.1 Naive

$$\text{Episode}_i(E_i) \equiv S_0^i \rightarrow A_0^i \rightarrow R_1^i \rightarrow S_1^i \rightarrow A_1^i \rightarrow R_2^i \cdots \rightarrow S_{T_i}$$

Episode E_i terminates at T_i .

- Let us define G_i to be the reward of E_i . $G_i \equiv \sum_{k=0}^{T_i} \gamma^k R_{k+1}^i$
- Estimate the value of π starting from s as $\hat{v}_\pi(s) = \frac{1}{m} \sum_{i=1}^m G_i$.
- Show by chernoff bounds that this is an OK estimate. We can use Chernoff as $\{G_i\}$ are independent, since the episodes $\{E_i\}$ are independent.
- *First-visit MC*: Average returns for the first time s is visited in an episode
- *Every-visit MC*: Average returns for every time s is visited in an episode.

Both of these asymptotally converge to the correct v_π .

2.1.1 First visit monte carlo policy evaluation

Run π from a fixed state s_0 for m times. This gives us m episodes. The i th episode is E_i , which terminates at step T_i .

$G(s, E_i)$ is defined as the return of π in run E_i , starting from the time instant of the first appearance of s in E_i till the final state. If state s occurs at time t_s , then $G(s, E_i) \equiv \sum_{j=t_s}^{T_i} \gamma^{(j-t_s)} R_{j+1}^i$. We start from $j+1$ since we want all rewards *after* our state. Note that the reward R_j is the reward granted *before* transitioning to the state S_j .

The value of a state s under the policy π is defined as:

$$\hat{v}_\pi \equiv \frac{1}{m} \sum_{i=1}^m G(s, E_i)$$

2.1.2 Every visit monte carlo

Every time we visit a state, we are able to find the return starting from that state.

```
-- / discount factor
gamma = 0.9

-- / an episode is a list of states,
-- with a possible (action,reward) pair
-- generating the next state
type Episode = [(S, Maybe (A, R))]

-- / Note the use of ParallelListComprehension!
reward :: Episode -> R
reward es = sum $ [gamma^i * r | i <- [1..] | (_, Just(_, r)) <- es]

-- / Return all possible tails, in order of longest
-- to shortest subsequence.
-- > tails [1, 2, 3] = [[1, 2, 3], [2, 3], [3], []]
tails :: Episode -> [Episode]
tails [] = []
tails xs = xs:tails (tail xs)

-- / Find the longest subsequence with start state s0
-- and calculate its reward. This assumes that
-- tails returns the longer subsequences first
firstVisit :: State -> Episode -> R
firstVisit s0 episodes =
  case dropWhile (\e -> fst <$> headMaybe e /= Just s0) (tails episodes) of
    (e:_) -> reward e
    _ -> 0

-- / Find all subsequences with start state s0, and calculate
-- their rewards
everyVisit :: State -> Episode -> R
everyVisit s0 episodes = sum $ do
  e <- tails episodes
  return $ if fst <$> headMaybe e /= Just s0 then 0 else reward e
```

For each run E_i and state s_0 , let $G(s_0, E_i, j)$ be the return of π in the run E_i for the j th occurrence of s_0 in E_i . Let $N_i(s)$ be the number of times state s has occurred in episode E_i .

$$\hat{v}_\pi(s) = \frac{1}{\sum_{i=1}^m N_i(s)} \sum_{i=1}^m \sum_{j=1}^{N_i(s)} G(s, E_i, j)$$

Note that $G(s, E_i, j)$ for a fixed state s and E_i is a dependent variable for different j . That is, $G(s, E_i, 1)$ is dependent on $G(s, E_i, 0)$.

We also want $q_{\pi}(s, a)$: The expected return starting from state s , taking an action a , and then following the policy π . (evaluation uses v , updates involve q).

2.2 Soft policy

Since our policy is deterministic, we cannot explore all state-action pairs. Therefore, we make our policy softly non-deterministic, by allowing transitions to all states with probability ϵ . This allows us to explore all state-action pairs.

- On-policy: Use same policy to generate the episode and update the policy.
- Off-policy: Use some policy to generate the episode, and update a different policy.

$$\begin{aligned} \pi : S \times A &\rightarrow \mathbb{R} \\ \pi(s|a) &= \begin{cases} 1 - \epsilon + \frac{\epsilon}{|A|} & \text{if } a = a^* \\ \frac{\epsilon}{|A|} & \text{otherwise} \end{cases} \end{aligned}$$

2.3 On policy first-visit MC control, for ϵ -soft policy

```
Q :: State -> Action -> Prob
Q s a = random
```

```
Returns :: State -> Action -> Real
Returns s a = 0
```

```
pi :: State -> Action -> Prob
pi = an arbitrary epsilon soft policy
```

```
- generate a new episode using pi
- for each (s, a) in the episode,
  G <- return following the first occurrence of (s, a)
  append G to Returns(s, a)
  Q(s, a) <- average>Returns(s, a))

- for each s in the episode: A* <- arg max a Q(s, a)
- pi <- a new epsilon soft policy based on A*
```

Because the optimal policy for a finite MDP is deterministic, we choose to keep the policy slightly away from deterministic: This way, we get to explore the space, while still being optimal.

We will prove that this actually does improve the policy. Let the new policy be π_{k+1} , and the current policy be π_k .

$$\begin{aligned} q_{\pi}(s, \pi_{k+1}(a)) &= \sum_{a \in \mathcal{A}(s)} \pi_{k+1}(a|s) q_{\pi_k}(s, a) \\ &= (1 - \epsilon) \max_{a \in \mathcal{A}(s)} q_{\pi_k}(s, a) + \sum_{a \in \mathcal{A}(s)} \frac{\epsilon}{|\mathcal{A}(s)|} q_{\pi_k}(s, a) \\ &= (1 - \epsilon) q_{\pi_k}(s, a^*) + \sum_{a \in \mathcal{A}(s)} \frac{\epsilon}{|\mathcal{A}(s)|} q_{\pi_k}(s, a) \end{aligned}$$

Note that $\pi'(a) = \frac{\pi_k(a|s) - \frac{\epsilon}{|\mathcal{A}(s)|}}{1 - \epsilon} \geq 0$. Also note that $\sum_a \pi'(a) = 1$. Now, since $\text{avg} \leq \max$, $\sum_a \pi'(a|s) q_{\pi_k}(s, a) \leq q_{\pi_k}(s, a^*)$.

$$\begin{aligned} &= (1 - \epsilon) q_{\pi_k}(s, a^*) + \sum_{a \in \mathcal{A}(s)} \frac{\epsilon}{|\mathcal{A}(s)|} q_{\pi_k}(s, a) \\ &\geq (1 - \epsilon) \sum_a \pi'(a|s) q_{\pi_k}(s, a) + \sum_{a \in \mathcal{A}(s)} \frac{\epsilon}{|\mathcal{A}(s)|} q_{\pi_k}(s, a) \\ &= (1 - \epsilon) \left[\frac{\sum_a \left[\pi_k(a|s) - \frac{\epsilon}{|\mathcal{A}(s)|} \right]}{1 - \epsilon} \right] q_{\pi_k}(s, a) + \sum_{a \in \mathcal{A}(s)} \frac{\epsilon}{|\mathcal{A}(s)|} q_{\pi_k}(s, a) \\ &= \sum_a \pi_k(a|s) q_{\pi_k}(s, a) \quad \left(\text{Cancellation of } \sum_{a \in \mathcal{A}(s)} \frac{\epsilon}{|\mathcal{A}(s)|} q_{\pi_k}(s, a) \right) \\ &= v_{\pi_k}(s) \end{aligned}$$

2.4 Off policy

We can use another policy μ to generate data, and we estimate q_{π} based on μ . Here, π is called the target policy, and μ is called the behaviour policy.

We need μ to *cover* π : $\pi(a|s) > 0 \implies \mu(a|s) > 0$. That is, every action taken by π must have non-zero probability to be taken by μ . That is, we can only choose to take actions that were taken by μ , since we can only really learn from the things that μ has done.